

Database Management Systems


Lecture 3

Transactions. Concurrency Control

- locking *protocols*
 - Strict Two-Phase Locking
 - Two-Phase Locking
- *deadlocks*
 - prevention (Wait-die, Wound-wait)
 - detection (waits-for graph, timeout mechanism)
- the *phantom* problem
- *isolation levels*
 - READ UNCOMMITTED
 - READ COMMITTED
 - REPEATABLE READ
 - SERIALIZABLE

Strict Two-Phase Locking (*Strict 2PL*)


- * before a transaction can read / write an object, it must acquire a S / X lock on the object
- * all the locks held by a transaction are released when it completes execution



- the Strict 2PL protocol allows only serializable schedules (only schedules with acyclic precedence graphs are allowed by this protocol)

Strict Two-Phase Locking

- the interleaving below is not allowed by Strict 2PL:



->

Strict Two-Phase Locking

T_1
XLock(CM)
R(CM)
W(CM)
...

T_2

XLock(CM)
...

- T_1 acquires an X lock on object CM, reads and writes CM
 - T_1 is still in progress when T_2 requests a lock on the same object, CM
 - T_2 cannot acquire an exclusive lock on CM, since T_1 already holds a conflicting lock on this object
 - T_1 will release its lock on CM only when it completes execution (with *commit* or *abort*)
 - since it cannot grant T_2 the requested lock on CM, the DBMS suspends T_2
- in this example, we denote by $XLock(O)$ the action of the current transaction requesting an X lock on object O

Strict Two-Phase Locking

T_1

XLock(CM)

R(CM)

W(CM)

XLock(CS)

R(CS)

W(CS)

commit

T_2

XLock(CM)

R(CM)

W(CM)

XLock(CS)

R(CS)

W(CS)

commit

- T_1 continues execution
- when T_1 commits, it releases both locks (X lock on CM, X lock on CS)
- T_2 can now be granted an X lock on CM
- T_2 can now proceed

Strict Two-Phase Locking


- the interleaving below is allowed by Strict 2PL:

T_1	T_2
XLock(CM)	
R(CM)	
W(CM)	
	XLock(CT)
	R(CT)
	W(CT)
	commit
XLock(CS)	
R(CS)	
W(CS)	
commit	

- in this example, since T_1 and T_2 are operating on separate data objects (CM, CT, CS), they can concurrently obtain all requested locks (an interleaved execution where T_1 and T_2 are only reading the same data object is also allowed)

Two-Phase Locking (2PL)

- variant of Strict Two-Phase Locking
 - * before a transaction can read / write an object, it must acquire a S / X lock on the object
 - * once a transaction releases a lock, it cannot request other locks
- phase 1 - *growing phase*
 - transaction acquires locks
- phase 2 - *shrinking phase*
 - transaction releases locks



Two-Phase Locking

- C – set of transactions
- $Sch(C)$ – set of schedules for C
- if all transactions in C obey 2PL, then any schedule $S \in Sch(C)$ that completes normally is serializable

Two-Phase Locking

- the following execution is allowed by the protocol:
- T1 can release its X lock on D prior to completion
- so T2 can acquire an X lock on D while T1 is still in progress
- however, T1 cannot acquire any other locks once it releases a lock (in this case, its X lock on D)

T1
R(D)
W(D)

T2

R(D)
W(D)
R(F)
W(F)
commit

...
commit

Two-Phase Locking

T1	T2
R(D)	
W(D)	
	R(D)
	W(D)
	R(F)
	W(F)
	commit

t

- suppose T1 is forced to terminate at time t
 - undo T1's updates => T2's update of D is lost (i.e., partial effects, as T2 also changed the value of F)
- problem – T1 released its exclusive lock on D prior to completion (under Strict 2PL, T1 can release its lock on D, as well as any other locks, only when it commits / aborts)


Strict Schedules

- if transaction T_i has written object A , then transaction T_j can read and / or write A only after T_i 's completion (commit / abort)
- strict schedules:
 - avoid cascading aborts
 - are recoverable schedules
 - if a transaction is aborted, its operations can be undone
- Strict 2PL only allows strict schedules

Deadlocks

- lock-based concurrency control techniques can lead to deadlocks
- deadlock
 - cycle of transactions waiting for one another to release a locked resource
 - normal execution can no longer continue without an external intervention, i.e., deadlocked transactions cannot proceed until the deadlock is resolved
- deadlock management
 - deadlock prevention
 - deadlock detection
 - allow deadlocks to occur and resolve them when they arise

Deadlocks



```
T1
BEGIN TRAN
XLock(V)
Read(V)
V := V + 100
Write(V)
XLock(Y)
Wait
Wait
...
```

```
T2
BEGIN TRAN
XLock(Y)
Read(Y)
Y := Y * 5
Write(Y)
XLock(V)
Wait
Wait
...
```

- T1 cannot obtain an X lock on Y, since T2 holds a conflicting lock on Y
- similarly, T2 cannot obtain an X lock on V, since T1 holds a conflicting lock on V


Deadlocks - Prevention

- assign transactions timestamp-based priorities (each transaction has a timestamp - the moment it begins execution)
- the lower the timestamp, the older the transaction
- the older a transaction is, the higher its priority, with the oldest transaction having the highest priority


- 2 deadlock prevention policies: Wait-die and Wound-wait

Deadlocks - Prevention

- assume T_1 wants to access an object locked by T_2 (with a conflicting lock)
 - Wait-die
 - if T_1 's priority is higher, T_1 can wait; otherwise, T_1 is aborted
- in the following execution, 2 transactions are reading and / or writing 2 objects, x and y; T_1 's priority is higher:




Deadlocks - Prevention



- T1 requests an X lock on object y, which is already locked with a conflicting lock by T2
- since T1 has a higher priority, it is allowed to wait
- T2 asks for an X lock on object x, already locked with a conflicting lock by T1
- since T2 has a lower priority, it is aborted
- T1 now obtains the requested lock on object y and proceeds with the write operation

Deadlocks - Prevention

- assume T_1 wants to access an object locked by T_2 (with a conflicting lock)
 - Wound-wait
 - if T_1 's priority is higher, T_2 is aborted; otherwise, T_1 can wait



- T1 requests an X lock on object y, which is already locked with a conflicting lock by T2
- since T1 has a higher priority, T2 is aborted
- T1 obtains the requested lock on object y and continues execution

Deadlocks - Prevention

- under these policies (Wait-die / Wound-wait), deadlock cycles cannot develop
- if an aborted transaction is restarted, it's assigned its original timestamp



Deadlocks - Detection

a. waits-for graph

- structure maintained by the lock manager to detect deadlock cycles
 - a node / active transaction
 - arc from T_i to T_j if T_i is waiting for T_j to release a lock
- cycle in the graph => deadlock
- DBMS periodically checks whether there are cycles in the waits-for graph

Deadlocks - Detection



a. waits-for graph



Deadlocks - Detection

a. waits-for graph


- if operation $X(A)$ is also part of T_3 , there will be an arc from T_3 to T_1 in the graph (since T_1 holds a conflicting lock on A , and T_3 is waiting for T_1 to release this lock)
- the graph has a cycle (T_1 is also waiting for T_2 to release a lock, which in turn is waiting for T_3 to release a lock)
=> deadlock
- aborting a transaction that appears on a cycle allows several other transactions to proceed



Deadlocks - Detection

a. waits-for graph

T_1	T_2	T_3	T_4
S(A)	S(A)	S(A)	
R(A)	R(A)	R(A)	
			S(B)
			R(B)
X(B)	X(B)	X(B)	
			X(A)
...



Deadlocks – Choosing the Deadlock Victim

- possible criteria to consider when choosing the deadlock victim
 - the number of objects modified by the transaction
 - the number of objects that are to be modified by the transaction
 - the number of locks held
- the policy should be “fair”, i.e., if a transaction is repeatedly chosen as a victim, it should be eventually allowed to proceed

Deadlocks - Detection

b. timeout mechanism

- very simple, practical method of detecting deadlocks
- if a transaction T has been waiting too long for a lock on an object, a deadlock is assumed to exist and T is terminated

The Phantom Problem

example 1. Researchers[RID, ..., ImpactFactor, Age]

- Page1: <R1, 5, 30>, <R2, 5, 20>
- Page2: <R3, 5, 100>, <R4, 5, 90>
- Page3: <R8, 6, 18>, <R9, 6, 19>
- concurrent transactions T1 and T2
 - transaction T1
 - retrieve the age of the oldest researcher for each of the impact factor values 5 and 6
 - transaction T2
 - add a new researcher with impact factor 5
 - remove researcher R9
- T1 and T2 obey Strict 2PL

The Phantom Problem

- Page1: <R1, 5, 30>, <R2, 5, 20>
- Page2: <R3, 5, 100>, <R4, 5, 90>
- Page3: <R8, 6, 18>, <R9, 6, 19>

- T1 identifies and locks pages holding researchers with IF 5 (Page1, Page2)
- T1 computes max age for IF 5 (100)
- T2 acquires X locks on: Page4 (onto which it adds a new researcher with IF 5 and age 102) and Page3 (from which it deletes researcher R9)
- T2 then commits, releasing all its locks
- T1 now obtains an S lock on Page3 (containing all researchers with IF 6), and computes max age for IF 6 (18)

T1

SLock(Page1)

SLock(Page2)

compute max age for IF 5 => 100

SLock(Page3)

compute max age for IF = 6 => 18

...

T2

XLock(Page4)

XLock(Page3)

add record <R5, 5, 102> to Page4

delete researcher R9

commit – all locks are released

The Phantom Problem

- Page1: <R1, 5, 30>, <R2, 5, 20>
- Page2: <R3, 5, 100>, <R4, 5, 90>
- Page3: <R8, 6, 18>, <R9, 6, 19>
- outcome of interleaved schedule on the right:
 - IF 5, Max Age 100
 - IF 6, Max Age 18
- outcome of serial schedule (T1T2):
 - IF 5, Max Age 100
 - IF 6, Max Age 19
- outcome of serial schedule (T2T1):
 - IF 5, Max Age 102
 - IF 6, Max Age 18

T1	T2
SLock(Page1)	
SLock(Page2)	
compute max age for IF 5 => 100	
	XLock(Page4)
	XLock(Page3)
	add record <R5, 5, 102> to Page4
	delete researcher R9
	commit – all locks are released
SLock(Page3)	
compute max age for IF = 6 => 18	
...	

->

The Phantom Problem

=> the interleaved schedule is not serializable (no serial schedule over the same set of transactions has the same outcome)

- however, the schedule is conflict serializable (the precedence graph is acyclic)

=> in the presence of insert operations, i.e., if new objects can be added to the database, conflict serializability does not guarantee serializability

The Phantom Problem example 2.

T1

```
SELECT *  
FROM Students  
WHERE GPA >= 8
```

```
SELECT *  
FROM Students  
WHERE GPA >= 8
```

...

T2

```
INSERT INTO Students VALUES  
(12, 'Mara', 'Dobse', 10)  
COMMIT
```

result set for T1's query

row corresponding to student with sid 12
is not in the result set

row corresponding to student with sid 12
now appears in the result set

Transaction Support in SQL - Isolation Levels

- *isolation level*
 - determines the degree to which a transaction is exposed to the operations of other concurrently running transactions
- greater concurrency -> concurrency anomalies

Transaction Support in SQL - Isolation Levels

- isolation levels
 - READ UNCOMMITTED
 - READ COMMITTED
 - REPEATABLE READ
 - SERIALIZABLE
- isolation levels can be set with the following command:
 - SET TRANSACTION ISOLATION LEVEL *isolevel*
 - *e.g.*, SET TRANSACTION ISOLATION LEVEL REPEATABLE READ
- *dirty writes* are not allowed under any isolation level (*dirty write* - a transaction T1 modifies an object previously written by an ongoing transaction T2)

Isolation Levels

- READ UNCOMMITTED

- a transaction T can read data modified by an ongoing transaction (uncommitted data)
- lowest degree of isolation
- no S locks when reading data
- dirty reads ✓
- unrepeatable reads ✓
- phantoms ✓

Isolation Levels

- READ COMMITTED

- a transaction T can only read committed data
- however, an object read by T can be changed by another transaction while T is in progress
- a transaction must acquire an exclusive lock prior to writing an object
- a transaction must acquire a shared lock prior to reading an object (i.e., the last transaction that modified the object is finished)
- exclusive locks are released at the end of the transaction
- shared locks are immediately released
- dirty reads ✗
- unrepeatable reads ✓
- phantoms ✓

Isolation Levels

- REPEATABLE READ

- a transaction T can only read committed data
- no object read by T can be changed by another transaction while T is in progress (if T reads an object O twice, no transaction can modify O between T's reads)
- a transaction must acquire an exclusive lock prior to writing an object
- a transaction must acquire a shared lock prior to reading an object
- exclusive locks are released at the end of the transaction
- shared locks are released at the end of the transaction
- dirty reads ✗
- unrepeatable reads ✗
- phantoms ✓

Isolation Levels

- SERIALIZABLE

- a transaction T can only read committed data
- no object read by T can be changed by another transaction while T is in progress
- if T reads a set of objects based on a search predicate, this set cannot be changed by other transactions while T is in progress
- a transaction must acquire locks on objects before reading / writing them
- a transaction also acquires locks on sets of objects that must remain unmodified
 - if query `SELECT * FROM Students WHERE GPA >= 8` is executed twice within a transaction, it must return the same answer set
- locks are held until the end of the transaction

Isolation Levels

- SERIALIZABLE
 - highest degree of isolation
 - dirty reads ✕
 - unrepeatable reads ✕
 - phantoms ✕

* Jim Gray *

References

- [Ra02] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (3rd Edition), McGraw-Hill, 2002
- [Le99] LEVENE, M., LOIZOU, G., A Guided Tour of Relational Databases and Beyond, Springer, 1999
- [Ga09] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book (2nd Edition), Pearson, 2009
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si19] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts (7th Edition), McGraw-Hill, 2019
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems,
<http://infolab.stanford.edu/~ullman/fcdb.html>