

Database Management Systems

Lecture 9

Evaluating Relational Operators

Query Optimization

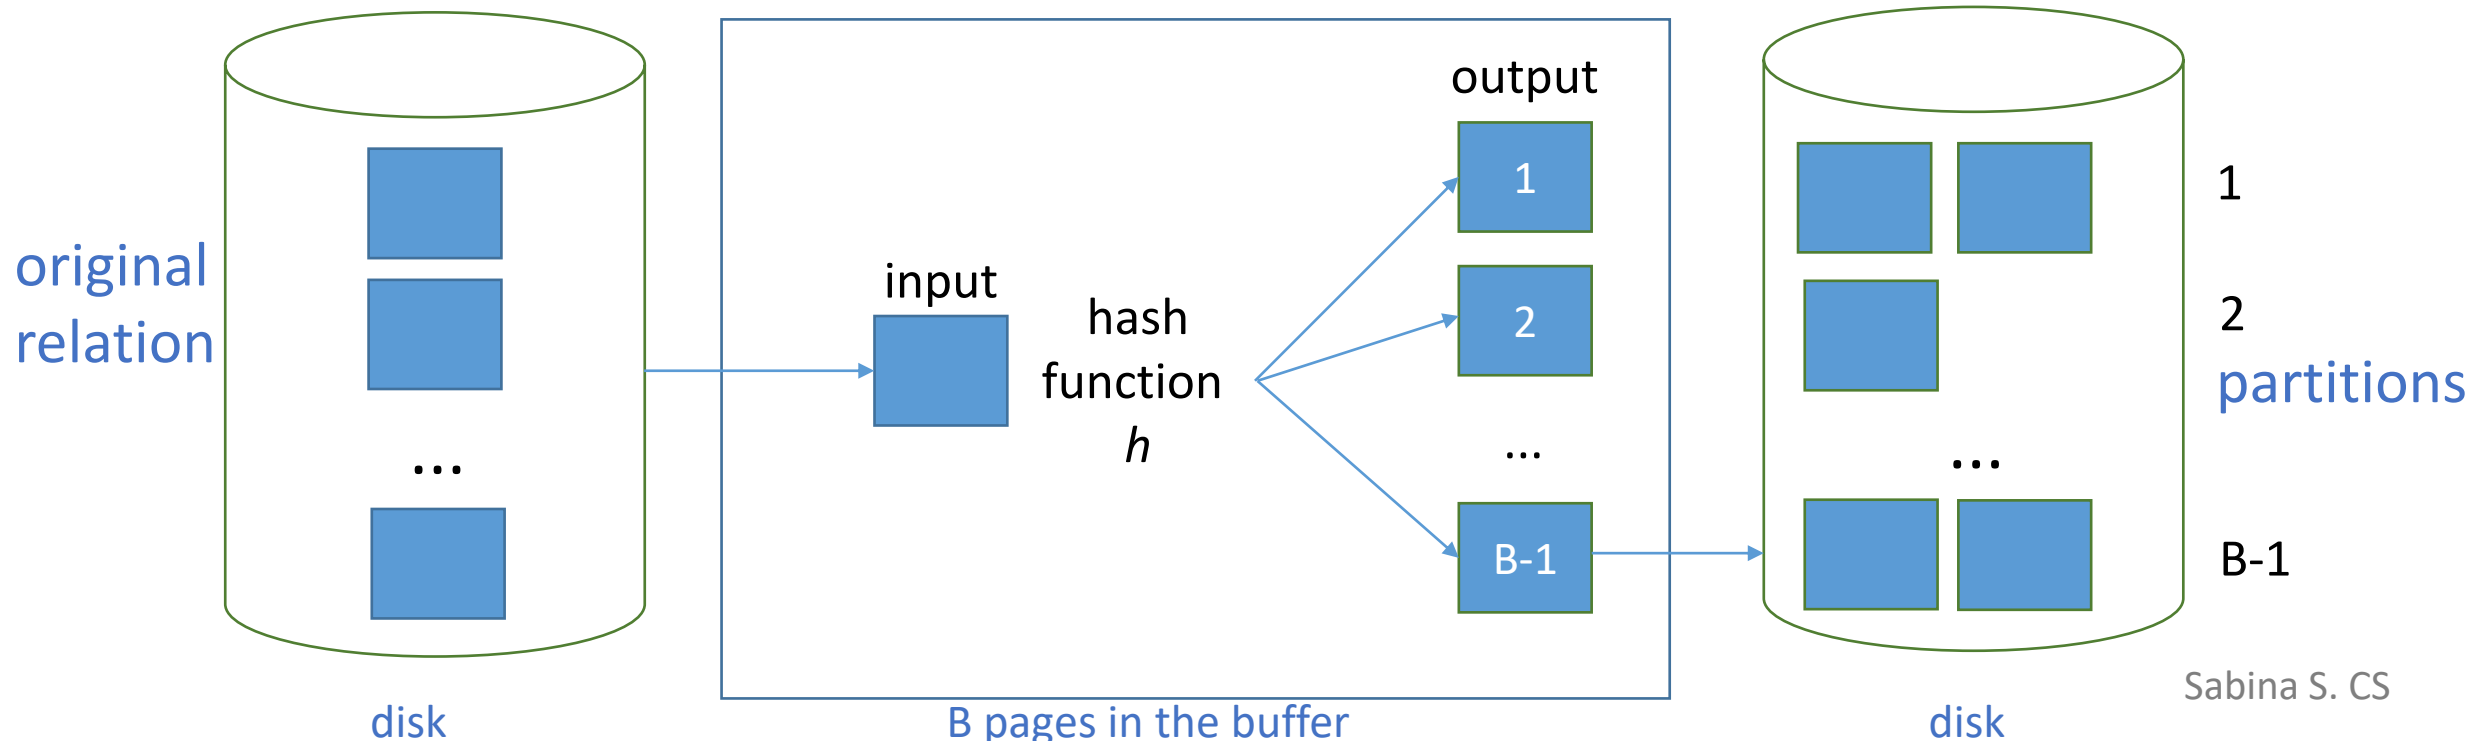
- running example - schema
 - Students (SID: integer, SName: string, Age: integer)
 - Courses (CID: integer, CName: string, Description: string)
 - Exams (SID: integer, CID: integer, EDate: date, Grade: integer, FacultyMember: string)
- Students
 - every record has 50 bytes
 - there are 80 records / page
 - 500 pages of Students tuples
- Courses
 - every record has 50 bytes
 - there are 80 records / page
 - 100 pages of Courses tuples
- Exams
 - every record has 40 bytes
 - there are 100 records / page
 - 1000 pages of Exams tuples

Hash Join - equality join, one join column: $E \bowtie_{i=j} S$

- phases: partitioning (building phase) & probing (matching phase)
- partitioning phase:
 - there are B pages available in the buffer:
 - use one page as the input buffer page
 - and the remaining $B-1$ pages as output buffer pages
 - choose a hash function h that distributes tuples uniformly to one of $B-1$ partitions
 - hash E and S on the join column (the i^{th} column of E , the j^{th} column of S) with the same hash function h

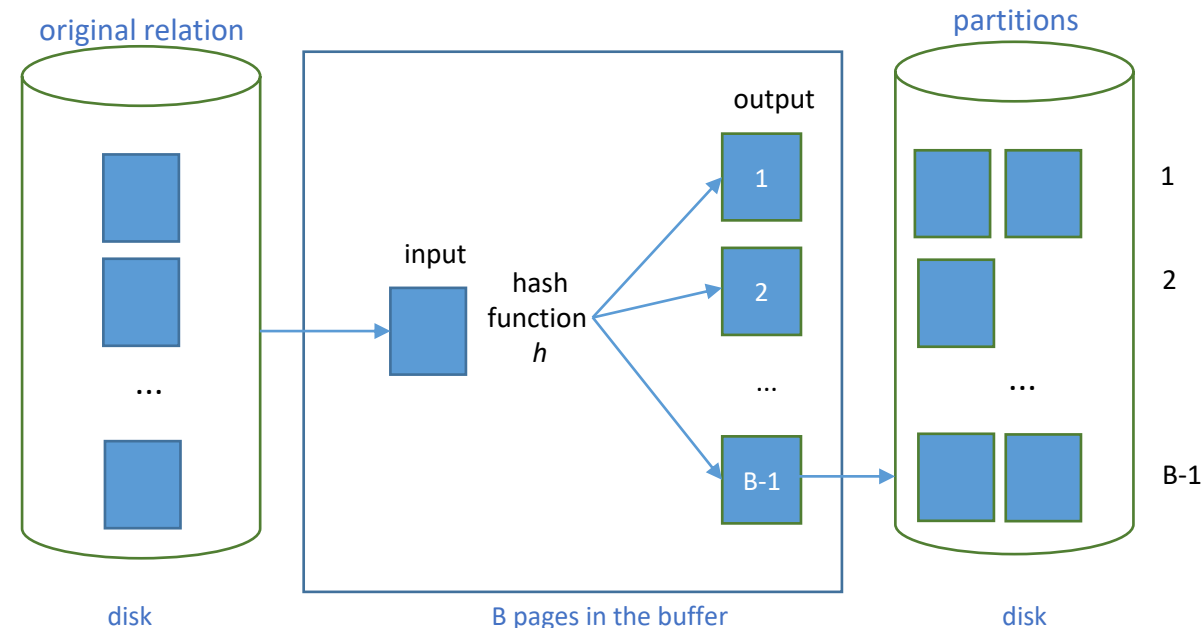
Hash Join

- hash E on the join column with hash function h (similarly for S):
 - for each tuple e in E, compute $h(e_i)$
(e_i : the value of the i^{th} column in tuple e)
 - add tuple e to the output buffer page that it is hashed to by h (buffer page $h(e_i)$)
 - when an output buffer page fills up, flush the page to disk



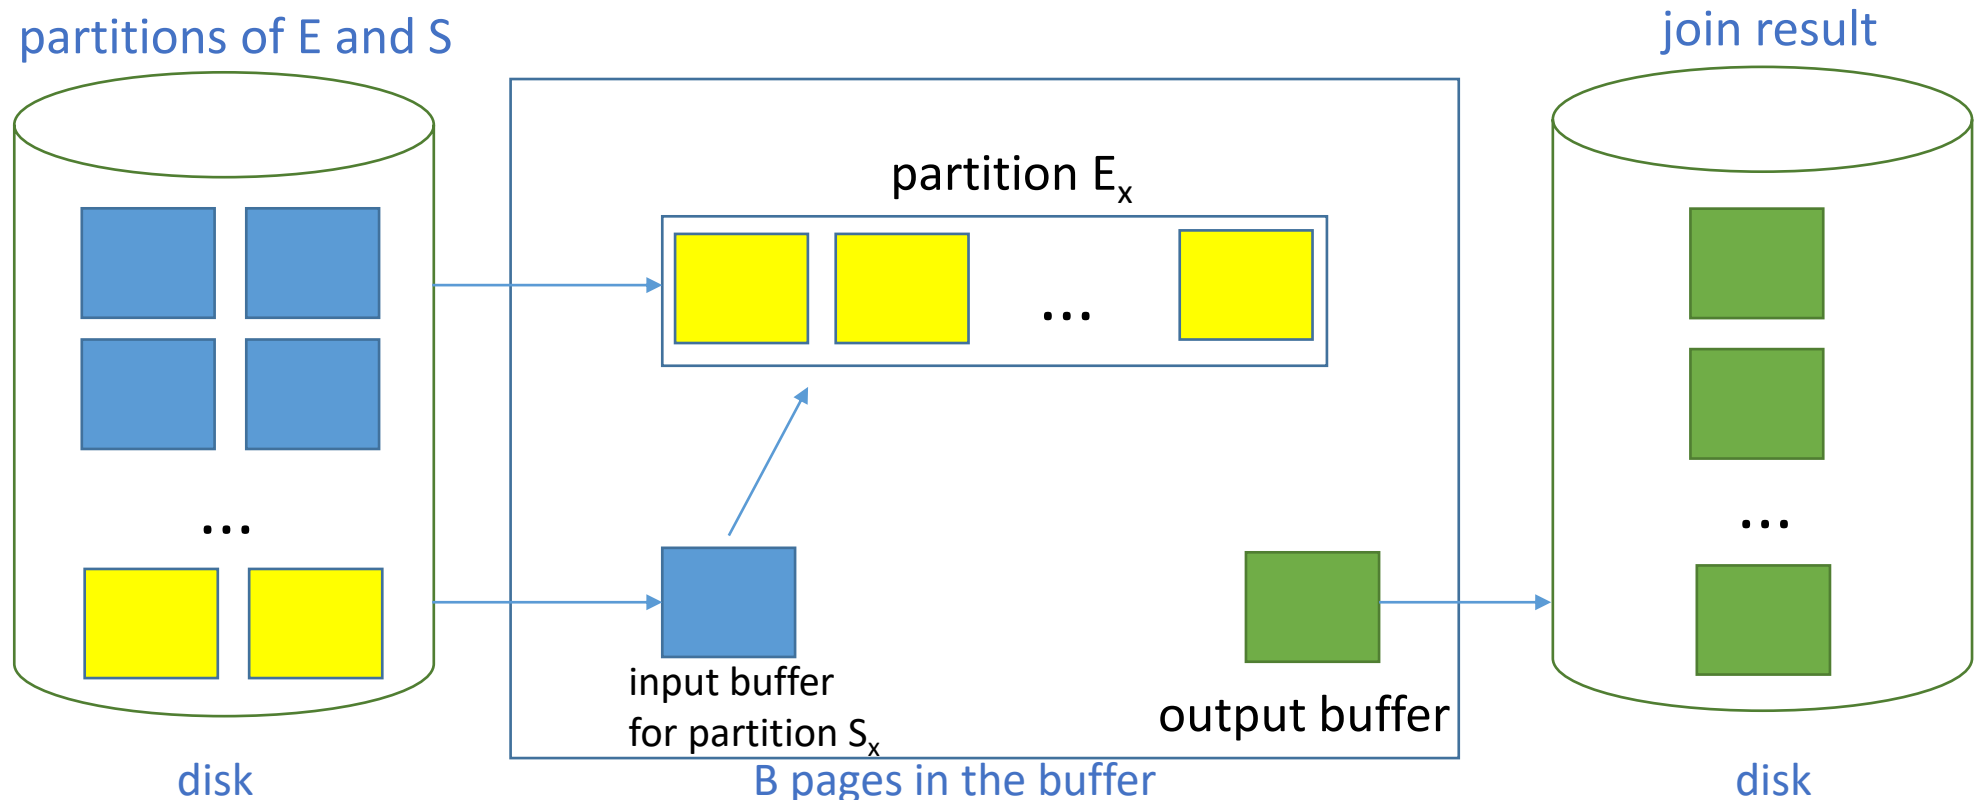
Hash Join

- partitioning phase \Rightarrow *partitions* of E (E_1, E_2 , etc) and S (S_1, S_2 , etc) on disk
- partition = collection of tuples that have the same hash value
- tuples in partition E_1 can only join with tuples in partition S_1 (they cannot join with tuples in partitions S_2 or S_3 , for instance, since these tuples have a different hash value)
- so to compute the join, we need to scan E and S only once (provided any partition of E fits in main memory)
- when reading in a partition E_k of E , we must scan only the corresponding partition S_k of S to find matching tuples (compare tuples e in E_k with tuples s in S_k to test the join condition *value of i^{th} column in E = value of j^{th} column in S*)



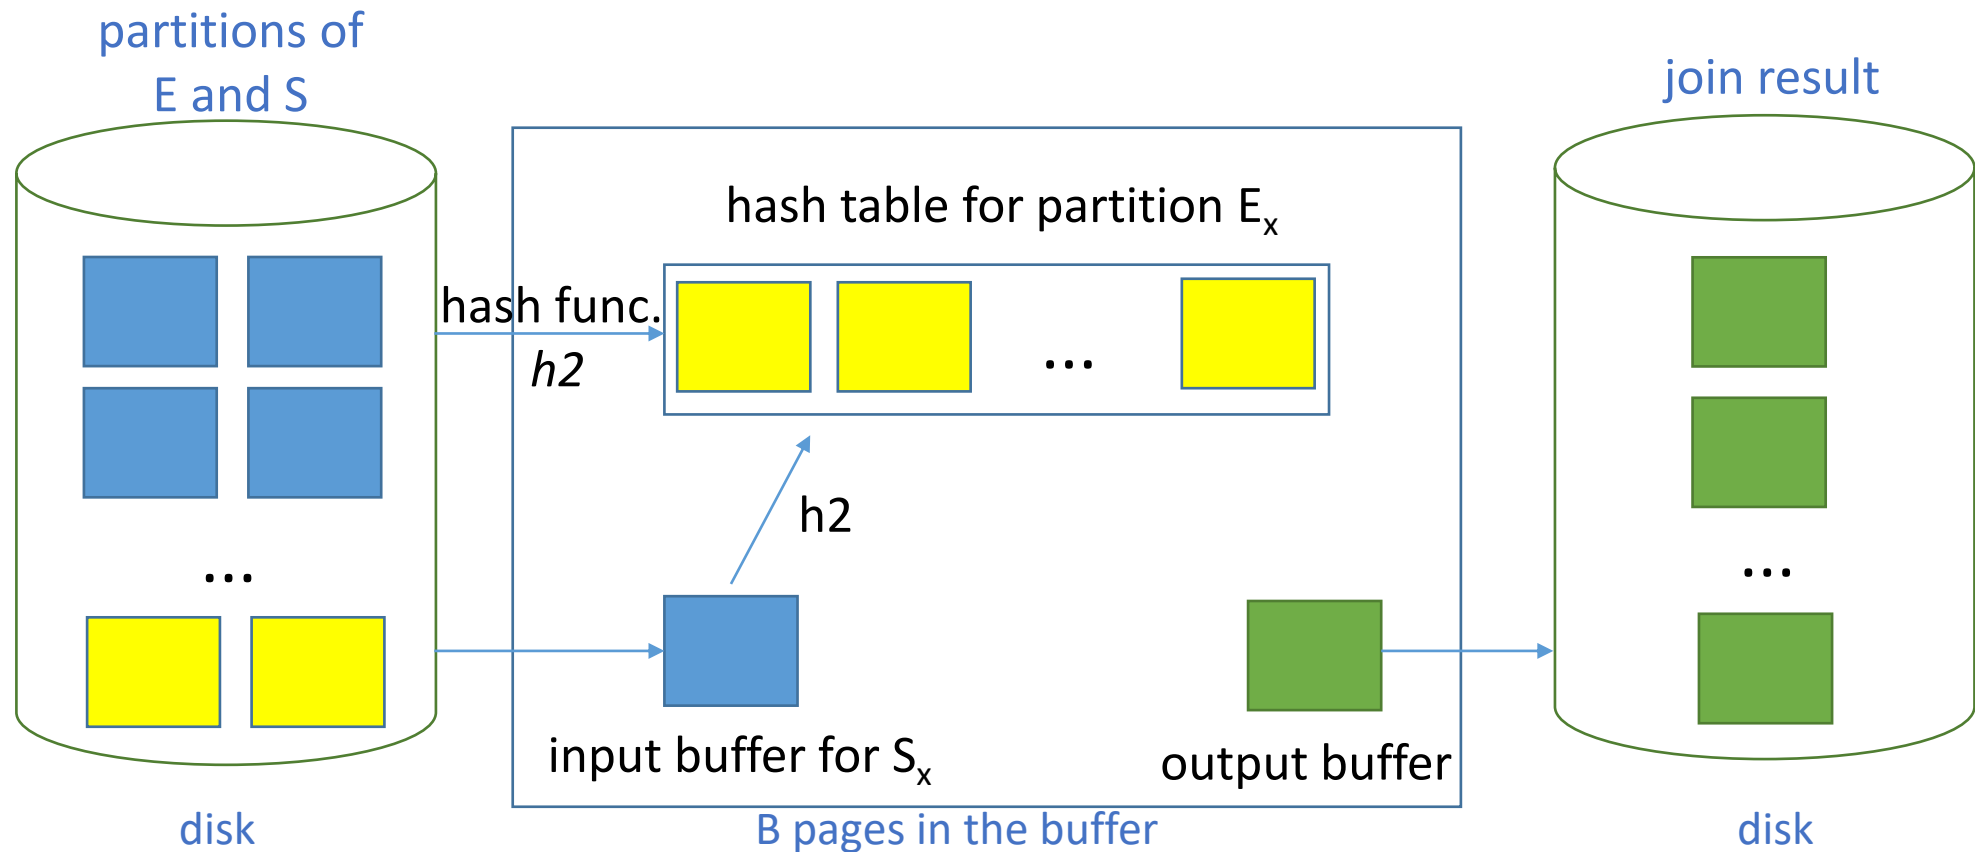
Hash Join

- probing phase:
 - read in a partition of the smaller relation (e.g., E) and scan the corresponding partition of S for matching tuples
 - use one page as the input buffer for S, one page as the output buffer, and the remaining pages to read in partitions of E



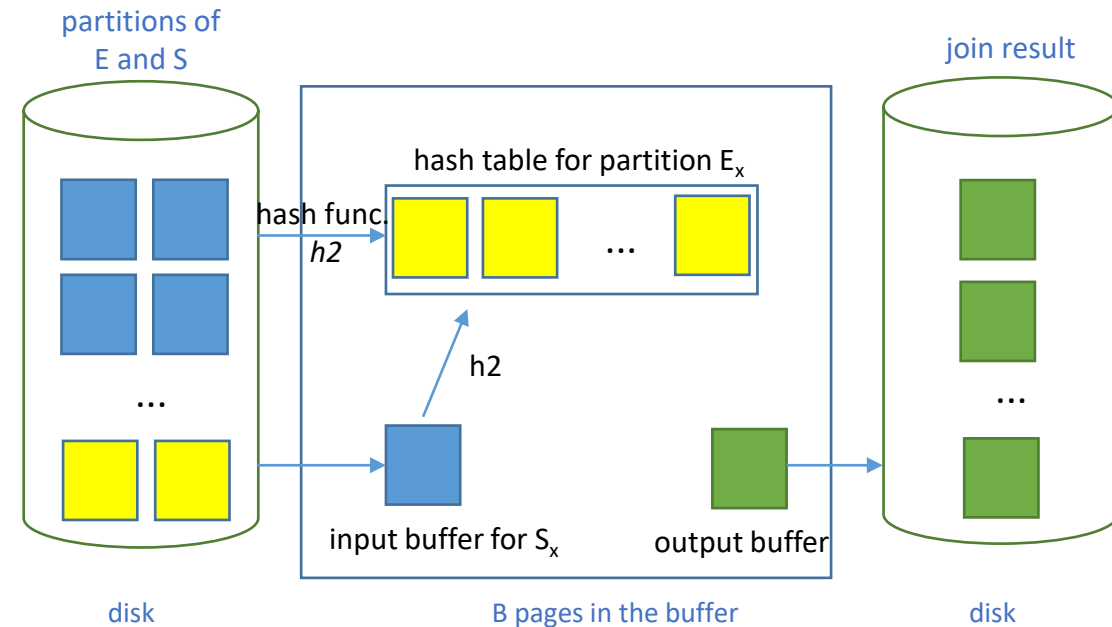
Hash Join

- probing phase:
 - in practice, to reduce CPU costs, an in-memory hash table is built, using a different function $h2$, for the E partition



Hash Join

- probing phase:
 - in practice, to reduce CPU costs, an in-memory hash table is built, using a different function $h2$, for the E partition
- consider a partition E_x of E
- build in-memory hash table for E_x using hash function $h2$ (the function is applied to the join column of E)
- for each tuple s in partition S_x , find matching tuples in the hash table using the hash value $h2(s_j)$
- result tuples $\langle e, s \rangle$ are written to output buffer
- once partitions E_x and S_x are processed, the hash table is emptied (to prepare for the next partition)



Hash Join

- cost:
 - partitioning:
 - both E and S are read and written once => cost: $2*(M+N)$ I/Os
 - probing:
 - scan each partition once => cost: $M+N$ I/Os
- => total cost: $3*(M+N)$ I/Os
- assumption: each partition fits into memory during probing
 - $3*(1000 + 500) = 4500$ I/Os

* E - M pages, p_E records / page *

* S - N pages, p_S records / page *

* 1000 pages * * 100 records / page *

* 500 pages * * 80 records / page *

Hash Join

- *partition overflow* – an E partition does not fit in memory during probing:
apply hash join technique recursively:
 - divide E, S into subpartitions
 - join subpartitions pairwise
 - if subpartitions don't fit in memory, apply hash join technique recursively

Hash Join

- memory requirements - objective: partition in E fits into main memory (S - similarly)
 - B buffer pages; need one input buffer => maximum number of partitions: B-1
 - size of largest partition: B - 2 (need one input buffer for S, one output buffer)
 - assume uniformly sized partitions => size of each E partition: $M/(B-1)$
=> $M/(B-1) < B-2$ => we need approximately $B > \sqrt{M}$
- if an in-memory hash table is used to speed up tuple matching => need a little more memory (because the hash table for a collection of tuples will be a little larger than the collection itself)

* E - M pages, p_E records / page *

* 1000 pages * * 100 records / page*

general join conditions

- equalities over several attributes
 - $E.SID = S.SID \text{ AND } E.attrE = S.attrS$
 - index nested loops join
 - Exams – inner relation:
 - build index on Exams with search key $\langle SID, attrE \rangle$ (if not already created)
 - can also use index on SID or index on attrE
 - Students – inner relation (similar)
 - sort-merge join
 - sort Exams on $\langle SID, attrE \rangle$, sort Students on $\langle SID, attrS \rangle$
 - hash join
 - partition Exams on $\langle SID, attrE \rangle$, partition Students on $\langle SID, attrS \rangle$
 - other join algorithms
 - essentially unaffected

general join conditions

- inequality comparison
 - $E.attrE < S.attrS$
 - index nested loops join
 - B+ tree index required
 - sort-merge join
 - not applied
 - hash join
 - not applied
 - other join algorithms
 - essentially unaffected

- * no join algorithm is uniformly superior to others
- choice of a good algorithm depends on:
 - size(s) of:
 - joined relations
 - buffer pool
 - available access methods

Selection

Q:

```
SELECT *  
FROM Exams E  
WHERE E.FacultyMember = 'Ionescu'
```

- use information in the selection condition to reduce the number of retrieved tuples
- e.g., $|Q| = 4$ (result set has 4 tuples), there's a B+ tree index on FacultyMember
 - it's expensive to scan E (1000 I/Os) to evaluate the query
 - should use the index instead
- selection algorithms based on the following techniques:
 - iteration, indexing

* E - M pages, p_E records / page * * 1000 pages * * 100 records / page*

Selection

- simple selections
 - $\sigma_{E.attr \text{ op } val}(E)$
- no index on *attr*, data not sorted on *attr*
 - must scan E and test the condition for each tuple
 - access path: file scan

=> cost: M I/Os = 1000 I/Os
- no index, sorted data (E physically sorted on *attr*)
 - binary search to locate 1st tuple that satisfies condition
and
 - scan E starting at this position until condition is no longer satisfied
 - access method: sorted file scan

Review lecture notes on *Relational Algebra, Indexes, DB – Physical Structure* (Databases course)

Selection

- simple selections
 - $\sigma_{E.attr \text{ op } val}(E)$
- no index, sorted data (E physically sorted on *attr*)
=> cost:
 - binary search: $O(\log_2 M)$
 - scan cost: varies from 0 to M
 - binary search on E
 - $\log_2 1000 \approx 10$ I/Os

Selection

- simple selections
 - $\sigma_{E.attr \text{ op } val}(E)$
- B+ tree index on *attr*
 - * search tree to find 1st index entry pointing to a qualifying E tuple
 - cost: typically 2, 3 I/Os
 - * scan leaf pages to retrieve all qualifying entries
 - cost: depends on the number of qualifying entries
 - * for each qualifying entry - retrieve corresponding tuple in E
 - cost: depends on the number of tuples and the nature of the index (clustered / unclustered)

Selection

- simple selections
 - $\sigma_{E.attr \text{ op } val}(E)$
- B+ tree index on *attr*
 - assumption
 - indexes use a2 or a3
 - a1-based index => data entry contains the data record => the cost of retrieving records = the cost of retrieving the data entries!
 - access path: B+ tree index
 - clustered index:
 - best access path when *op* is not *equality*
 - good access path when *op* is *equality*

Selection

- simple selections: $\sigma_{E.attr \text{ op } val}(E)$

- B+ tree index on *attr*

Q

```
SELECT *
```

```
FROM Exams E
```

```
WHERE E.FacultyMember < 'C%'
```

- names uniformly distributed with respect to 1st letter

=> $|Q| \approx 10,000$ tuples = 100 pages

- clustered B+ tree index on FacultyMember

=> cost of retrieving tuples: ≈ 100 I/Os (a few I/Os to get from root to leaf)

- non-clustered B+ tree index on FacultyMember

=> cost of retrieving tuples: up to 1 I/O per tuple (worst case) => up to 10.000 I/Os

* E - M pages, p_E records / page *

* 1000 pages * * 100 records / page*

Selection

- simple selections: $\sigma_{E.attr \text{ op } val}(E)$

- B+ tree index on *attr*

```
SELECT *
```

```
FROM Exams E
```

```
WHERE E.FacultyMember < 'C%'
```

- non-clustered B+ tree index on FacultyMember
 - refinement - sort rids in qualifying data entries by page-id
=> a page containing qualifying tuples is retrieved only once
 - cost of retrieving tuples: number of pages containing qualifying tuples (but such tuples are probably stored on more than 100 pages)
- range selections
 - non-clustered indexes can be expensive
 - could be less costly to scan the relation (in our example: 1000 I/Os)

Selection

- general selections
 - selections without disjunctions
- C - CNF condition without disjunctions
 - evaluation options:
 1. use the most selective access path
 - if it's an index I:
 - apply conjuncts in C that match I
 - apply rest of conjuncts to retrieved tuples
 - example
 - $c < 100 \text{ AND } a = 3 \text{ AND } b = 5$
 - can use a B+ tree index on c and check $a = 3 \text{ AND } b = 5$ for each retrieved tuple
 - can use a hash index on a and b and check $c < 100$ for each retrieved tuple

Selection

- general selections - selections without disjunctions
 - evaluation options:
 2. use several indexes - when several conjuncts match indexes using a2 / a3
 - compute sets of rids of candidate tuples using indexes
 - intersect sets of rids, retrieve corresponding tuples
 - apply remaining conjuncts (if any)
 - example: $c < 100 \text{ AND } a = 3 \text{ AND } b = 5$
 - use a B+ tree index on c to obtain rids of records that meet condition $c < 100$ (R_1)
 - use a hash index on a to retrieve rids of records that meet condition $a = 3$ (R_2)
 - compute $R_1 \cap R_2 = R_{int}$
 - retrieve records with rids in R_{int} (R)
 - check $b = 5$ for each record in R

Selection

- general selections
 - selections with disjunctions
- C - CNF condition with disjunctions, i.e., some conjunct J is a disjunction of terms
 - if some term T in J requires a file scan, testing J by itself requires a file scan
 - example: $a < 100 \vee b = 5$
 - hash index on b , hash index on c
 - => check both terms using a file scan (i.e., best access path: file scan)
- compare with the example below:
 - $(a < 100 \vee b = 5) \wedge c = 7$
 - hash index on b , hash index on c
 - => use index on c , apply $a < 100 \vee b = 5$ to each retrieved tuple (i.e., most selective access path: index)

Selection

- general selections
 - selections with disjunctions
- C - CNF condition with disjunctions
 - every term T in a disjunction matches an index

=> retrieve tuples using indexes, compute union

 - example
 - $a < 100 \vee b = 5$
 - B+ tree indexes on a and b
 - use index on a to retrieve records that meet condition $a < 100$ (R_1)
 - use index on b to retrieve records that meet condition $b = 5$ (R_2)
 - compute $R_1 \cup R_2 = R$
 - if all matching indexes use a2 or a3 => take union of rids, retrieve corresponding tuples

Projection

- $\Pi_{\text{SID, CID}}(\text{Exams})$

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- to implement projection:
 - eliminate:
 - unwanted columns
 - duplicates
- projection algorithms - *partitioning* technique:
 - sorting
 - hashing

Projection Based on Sorting

- step 1
 - scan $E \Rightarrow$ set of tuples containing only desired attributes (E')
 - cost:
 - scan E : M I/Os
 - write temporary relation E' : T I/Os
 - T depends on: number of columns and their sizes, T is $O(M)$
- step 2
 - sort tuples in E'
 - sort key: all columns
 - cost: $O(T \log T)$ (also $O(M \log M)$)
- step 3
 - scan sorted E' , compare adjacent tuples, eliminate duplicates
 - cost: T
- total cost: $O(M \log M)$

Projection Based on Sorting

* example

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- scan Exams: 1000 I/Os
- size of tuple in E': 10 bytes

=> cost of writing temporary relation E': 250 I/Os

- available buffer pages: 20
 - E' can be sorted in 2 passes
 - sorting cost: $2 * 2 * 250 = 1000$ I/Os
- final scan of E' - cost: 250 I/Os

=> total cost: $1000 + 250 + 1000 + 250 = 2500$ I/Os

* E – record size = 40 bytes * * 1000 pages * * 100 records / page*

Projection Based on Sorting

* example

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- scan Exams: 1000 I/Os
- size of tuple in E': 10 bytes

=> cost of writing temporary relation E': 250 I/Os

- available buffer pages: 257
 - E' can be sorted in 1 pass
 - sorting cost: $2 * 1 * 250 = 500$ I/Os
- final scan of E' - cost: 250 I/Os

=> total cost: $1000 + 250 + 500 + 250 = 2000$ I/Os

* E – record size = 40 bytes * * 1000 pages * * 100 records / page*

Projection Based on Sorting

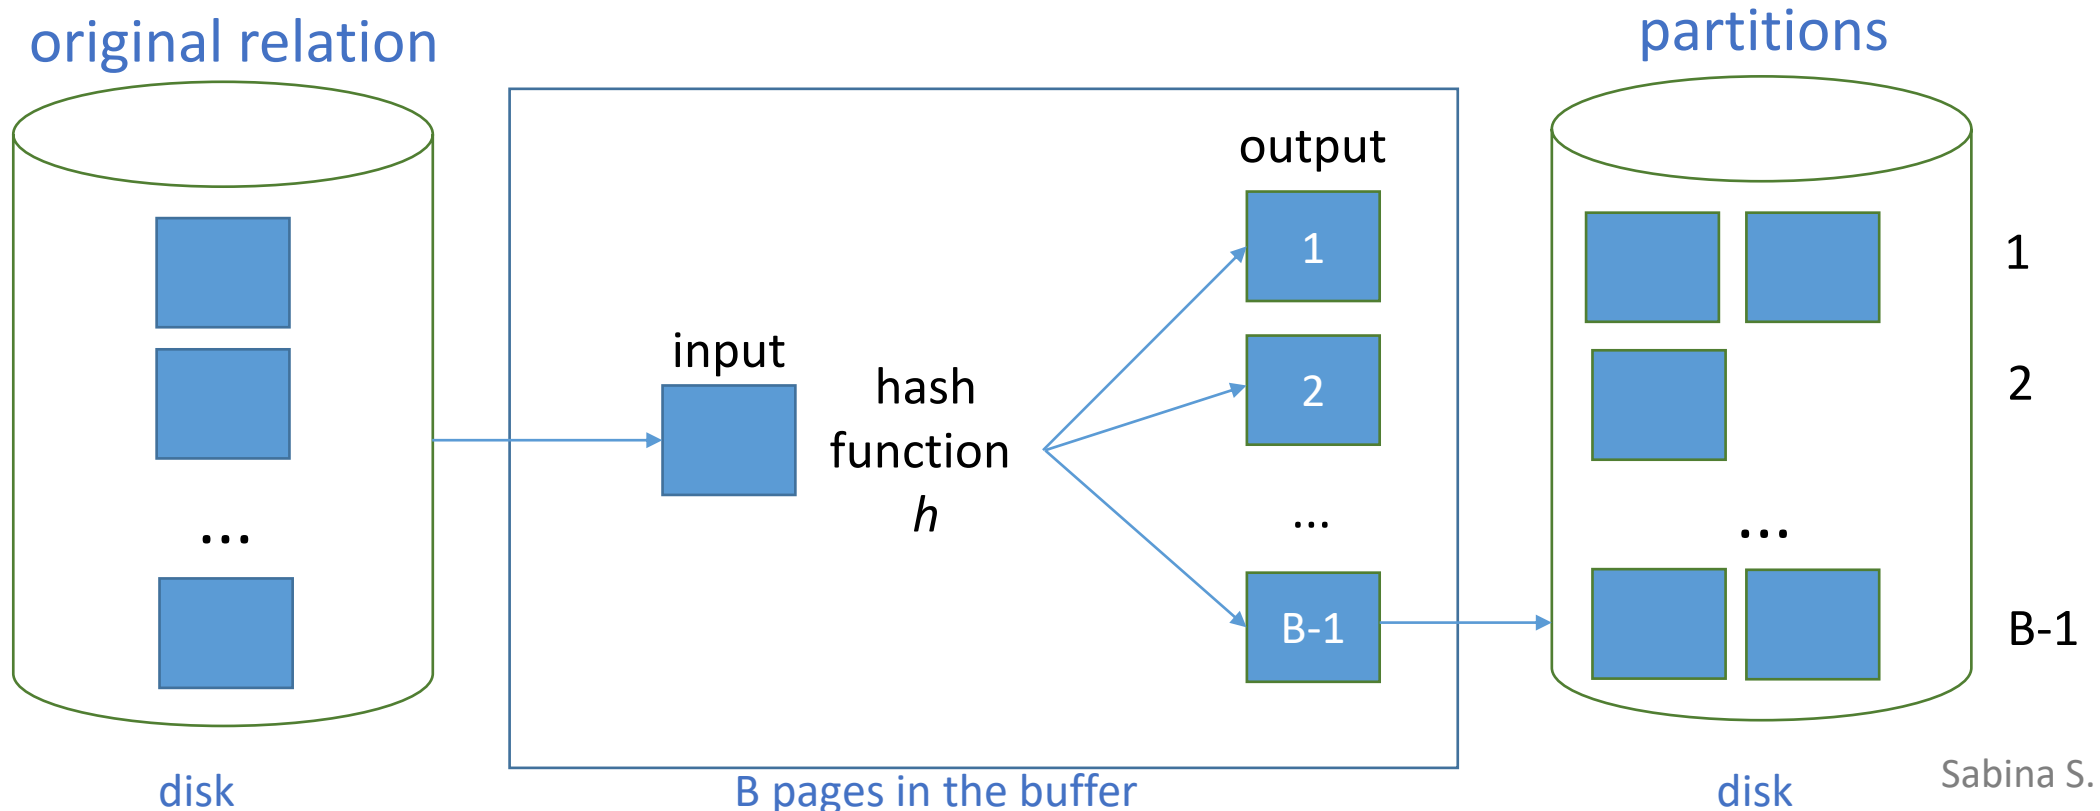
- improvement
 - adapt the sorting algorithm to do projection with duplicate elimination
 - modify pass 0 of External Merge Sort: eliminate unwanted columns
 - read in B pages from E
 - write out $(T/M) * B$ internally sorted pages of E'
 - refinement: write out $2*B$ internally sorted pages of E' (on average)
 - tuples in runs - smaller than input tuples
 - modify merging passes: eliminate duplicates
 - number of result tuples is smaller than number of input tuples

Projection Based on Sorting

- improvement
 - * example
 - pass 0:
 - scan Exams: 1000 I/Os
 - write out 250 pages:
 - 20 available buffer pages
 - 250 pages => 7 sorted runs about 40 pages long (except the last one, which is about 10 pages long)
 - pass 1:
 - read in all runs – cost: 250 I/Os
 - merge runs
 - total cost : $1000 + 250 + 250 = 1500$ I/Os

Projection Based on Hashing

- phases: partitioning & duplicate elimination
- partitioning phase:
 - 1 input buffer page – read in the relation one page at a time
 - hash function h – distribute tuples uniformly to one of $B-1$ partitions
 - $B-1$ output buffer pages – one output page / partition



Projection Based on Hashing

- partitioning phase:
 - read the relation using the input buffer page
 - for each tuple t :
 - discard unwanted fields \Rightarrow tuple t'
 - apply hash function h to t'
 - write t' to the output buffer page that it is hashed to by h

\Rightarrow B-1 partitions

- partition:
 - collection of tuples with:
 - common hash value
 - no unwanted fields
- tuples in different partitions are guaranteed to be distinct

Projection Based on Hashing

- duplicate elimination phase:
 - process all partitions:
 - read in partition P, one page at a time
 - build in-memory hash table with hash function $h_2 (\neq h)$ on all fields:
 - if a new tuple hashes to the same value as an existing tuple, compare them to check if they are distinct
 - eliminate duplicates
 - write duplicate-free hash table to result file
 - clear in-memory hash table
 - partition overflow
 - apply hash-based projection technique recursively (subpartitions)

Projection Based on Hashing

- cost
 - partitioning:
 - read E: M I/Os
 - write E': T I/Os
 - duplicate elimination:
 - read in partitions: T I/Os

=> total cost: $M + 2 * T$ I/Os
- Exams:
 - $1000 + 2 * 250 = 1500$ I/Os

Set Operations

- intersection, cross-product
 - special cases of join (join condition for intersection - equality on all fields, no join condition for cross-product)
- union, set-difference
 - similar
- union: $R \cup S$
 - sorting
 - sort R and S on all attributes
 - scan the sorted relations in parallel; merge them, eliminating duplicates
 - refinement
 - produce sorted runs of R and S, merge runs in parallel

Set Operations

- union: $R \cup S$
 - hashing
 - partition R and S with the same hash function h
 - for each S -partition
 - build in-memory hash table (using h_2) for the S -partition
 - scan corresponding R -partition, add tuples to hash table, discard duplicates
 - write out hash table
 - clear hash table

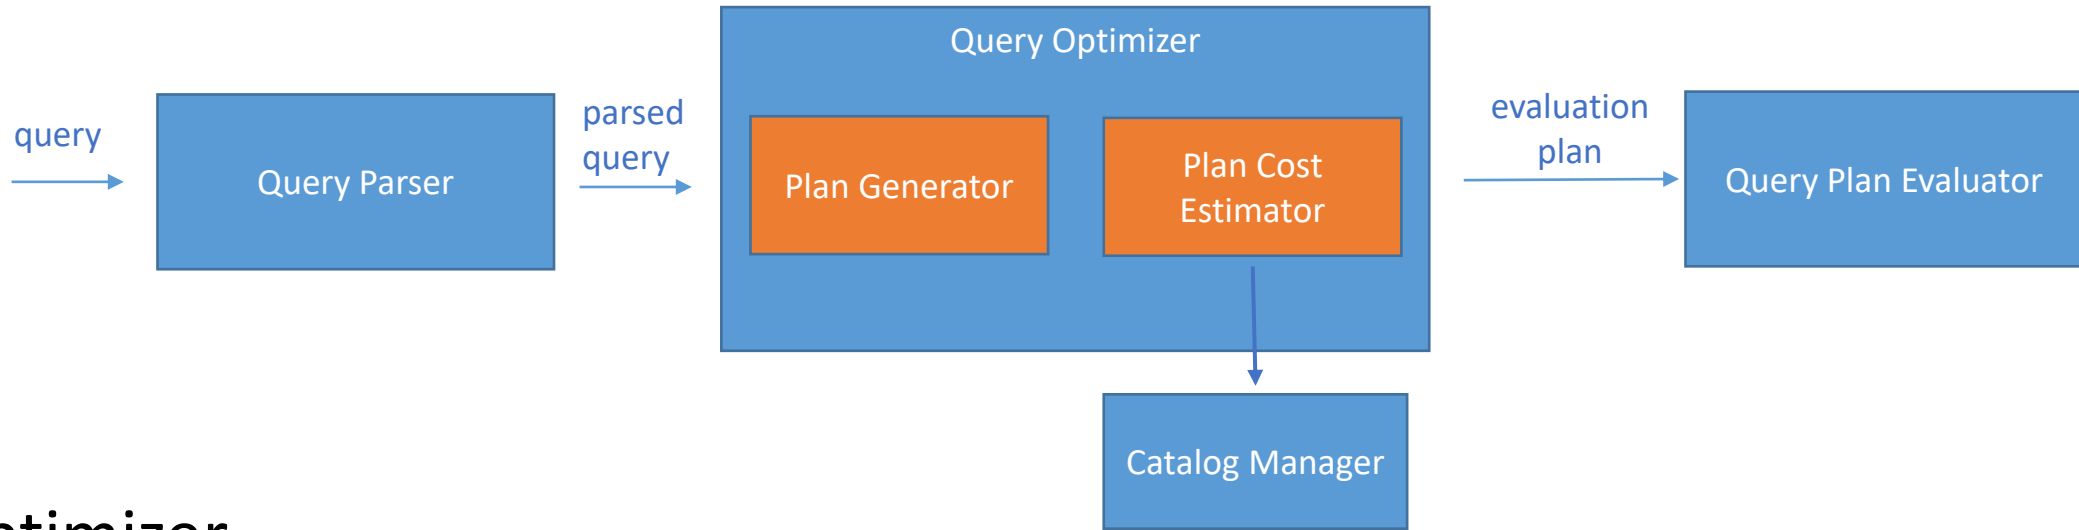
Aggregate Operations

- without grouping
 - scan relation
 - maintain *running information* about scanned tuples
 - COUNT - count of values retrieved
 - SUM - *total* of values retrieved
 - AVG - $\langle total, count \rangle$ of values retrieved
 - MIN, MAX - smallest / largest value retrieved
- with grouping
 - sort relation on the grouping attributes
 - scan relation to compute aggregate operations for each group
 - improvement: combine sorting with aggregation computation
 - alternative approach based on hashing

Aggregate Operations

- using existing indexes
 - index with a search key that includes all the attributes required by the query
 - work with the data entries in the index (instead of the data records)
 - attribute list in the GROUP BY clause is a prefix of the index search key (tree index)
 - get data entries (and records, if necessary) in the required order (i.e., avoid sorting)

Query Optimization



- optimizer
 - objective
 - given a query Q , find a good evaluation plan for a Q
 - generates alternative plans for Q , estimates their costs, and chooses the one with the least estimated cost
 - uses information from the system catalogs

- running example - schema
 - Students (SID: integer, SName: string, Age: integer)
 - Courses (CID: integer, CName: string, Description: string)
 - Exams (SID: integer, CID: integer, EDate: date, Grade: integer)
- Students
 - every record has 50 bytes
 - there are 80 records / page
 - 500 pages
- Courses
 - every record has 40 bytes
 - there are 100 records / page
 - 1 page
- Exams
 - every record has 40 bytes
 - there are 100 records / page
 - 1000 pages


Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 23
```

query

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

relational algebra expression

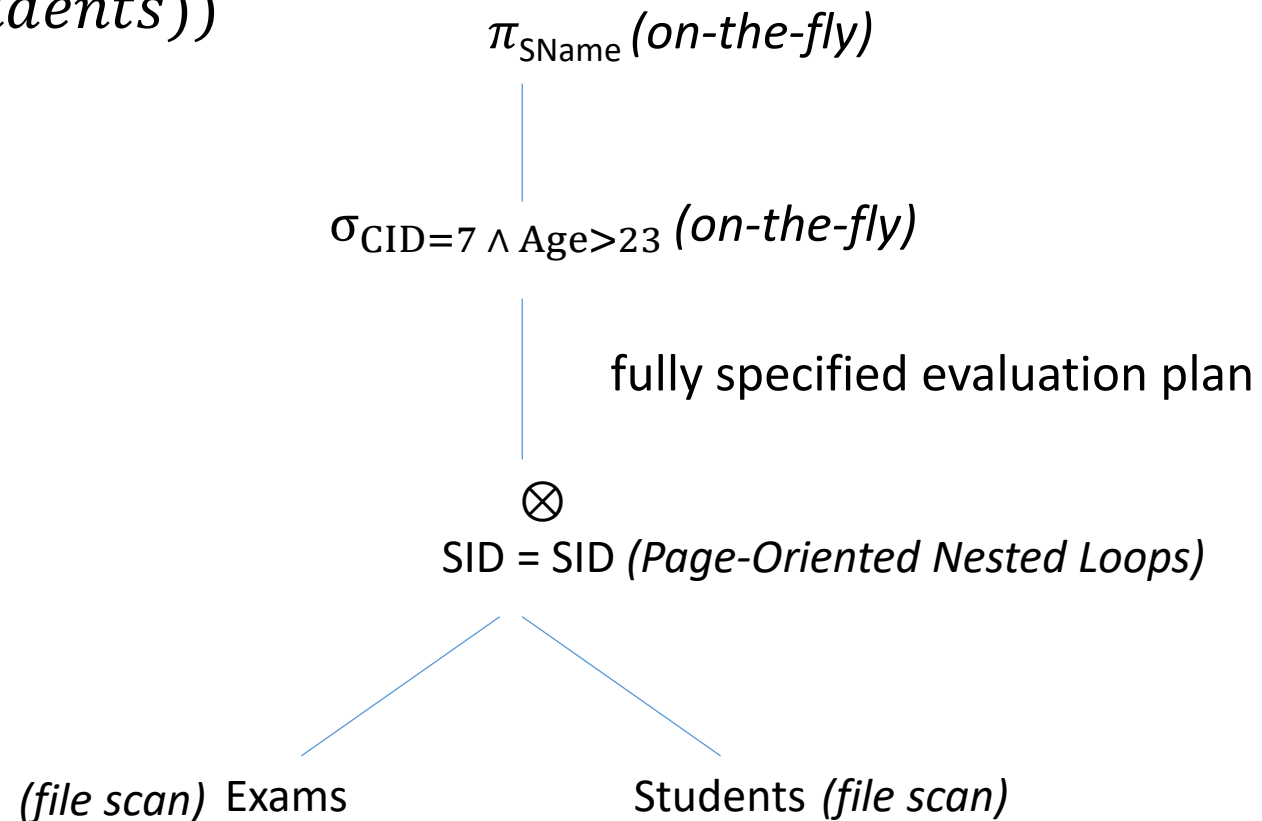


Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 23
```

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

- query evaluation plan
 - extended relational algebra tree
 - node – annotations
 - relation
 - access method
 - relational operator
 - implementation method

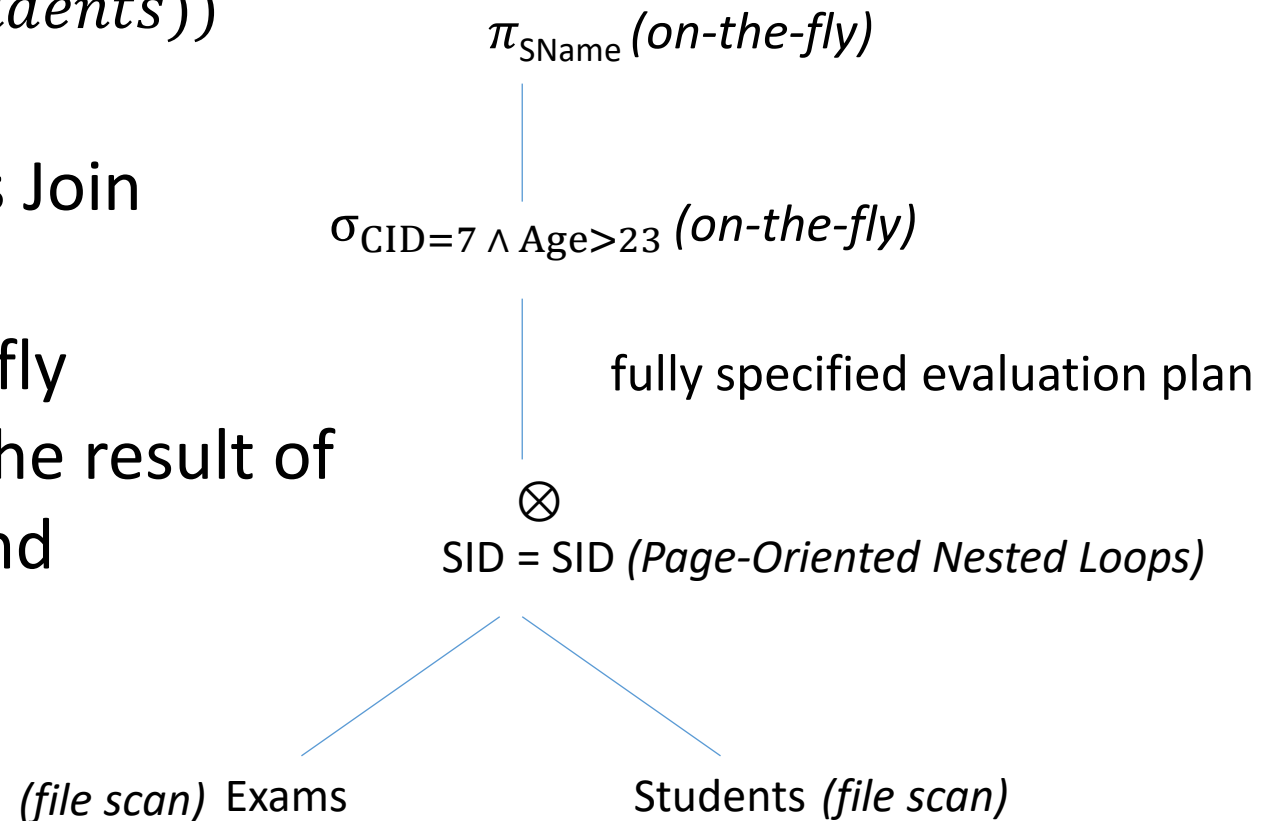


Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 23
```

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

- page-oriented Simplified Nested Loops Join
 - Exams – outer relation
- selection, projection applied on-the-fly to each tuple in the join result, i.e., the result of the join (before applying selection and projection) is not stored



Pipelined Evaluation

SELECT *

FROM Exams

WHERE $\frac{EDate > '1-1-2020'}{T1}$ AND $\frac{Grade > 8}{T2}$

$$\sigma_{Grade>8}(\sigma_{EDate>'1-1-2020'}(Exams))$$

- index I matches $T1$
- $v1$ - *materialization*
 - evaluate $T1$
 - write out result tuples to temporary relation R , i.e., tuples are *materialized*
 - apply the 2nd selection to R
 - cost: read and write R

Pipelined Evaluation

SELECT *

FROM Exams

WHERE $\frac{\text{EDate} > '1-1-2020'}{T1}$ AND $\frac{\text{Grade} > 8}{T2}$

- *v2 – pipelined evaluation*
 - apply the 2nd selection to each tuple in the result of the 1st selection as it is produced
 - i.e., 2nd selection operator is applied *on-the-fly*
 - saves the cost of writing out / reading in the temporary relation *R*

Query Blocks – Units of Optimization

- parse $Q \Rightarrow$ collection of query *blocks* \rightarrow passed on to the optimizer
- optimizer:
 - optimize one block at a time
- query *block* - SQL query:
 - without nesting
 - with exactly: one SELECT clause, one FROM clause
 - with at most: one WHERE clause, one GROUP BY clause, one HAVING clause
 - WHERE condition - CNF

Query Blocks – Units of Optimization

- query Q:

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = (SELECT MAX(S2.Age)
               FROM Students S2)
GROUP BY S.SID
HAVING COUNT(*) > 2
```

nested block

- decompose query into a collection of blocks without nesting

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```


Query Blocks – Units of Optimization

* block optimization

- express query block as a relational algebra expression

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```

$$\pi_{S.SID, MIN(E.EDate)}(\text{HAVING}_{COUNT(*) > 2}(\text{GROUP BY}_{S.SID}(\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value_from_nested_block}(Students \times Exams \times Courses))))$$

- GROUP BY, HAVING – operators in the extended algebra used for plans
- argument list of projection can include aggregate operations

Query Blocks – Units of Optimization

- query Q treated as a $\sigma \pi \times$ algebra expression
- the remaining operations in Q are performed on the result of the $\sigma \pi \times$ expression

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```

$$\pi_{S.SID, E.EDate}(\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value_from_nested_block}(Students \times Exams \times Courses))$$

- attributes in GROUP BY, HAVING are added to the argument list of projection
- aggregate expressions in the argument list of projection are replaced by their argument attributes

Query Blocks – Units of Optimization

* block optimization

- find best plan P for the $\sigma \pi \times$ expression
- evaluate P => result set RS
- sort/hash RS => groups
- apply HAVING to eliminate some groups
- compute aggregate expressions in SELECT for each remaining group

$\pi_{S.SID, MIN(E.EDate)}(\$
 $HAVING_{COUNT(*) > 2}(\$
 $GROUP BY_{S.SID}(\$
 $\pi_{S.SID, E.EDate}(\$
 $\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value_from_nested_block}(\$
 $Students \times Exams \times Courses))))$

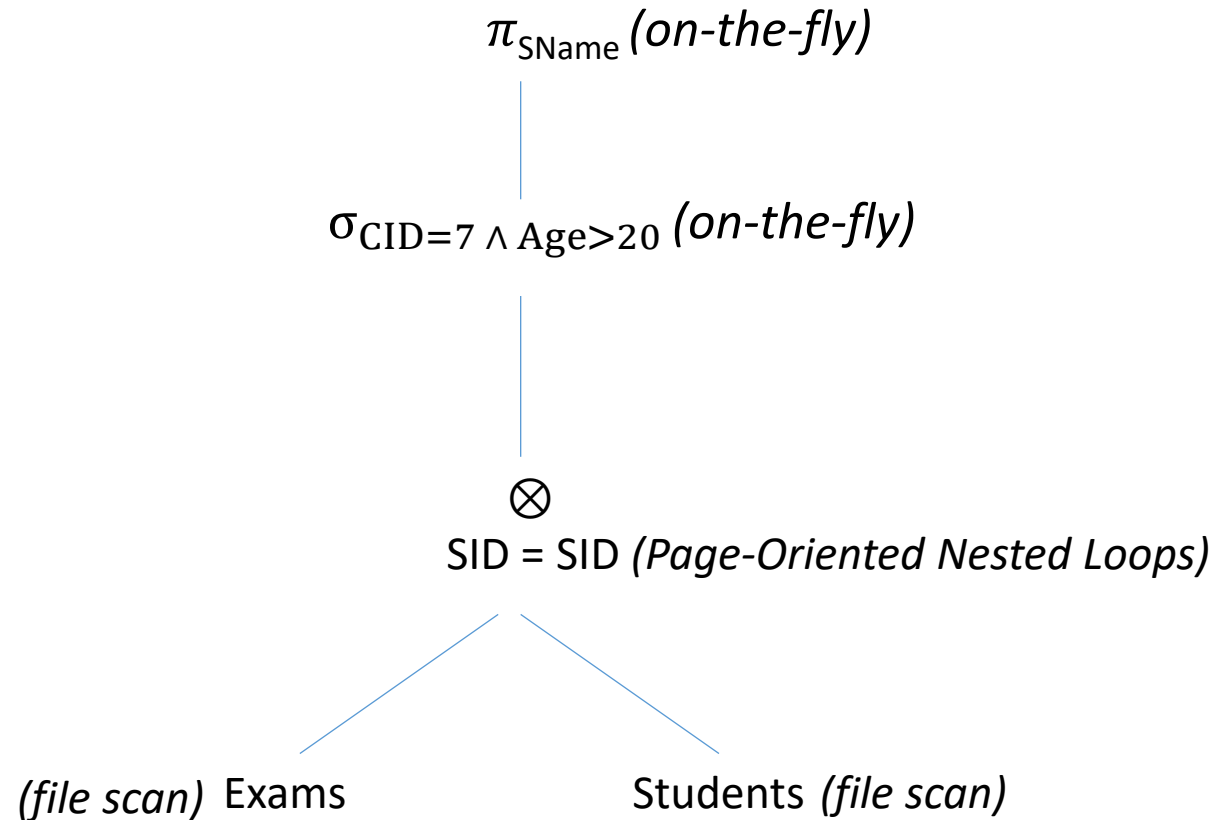
Motivating Example

* E - 1000 pages *

* S - 500 pages *

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 20
```

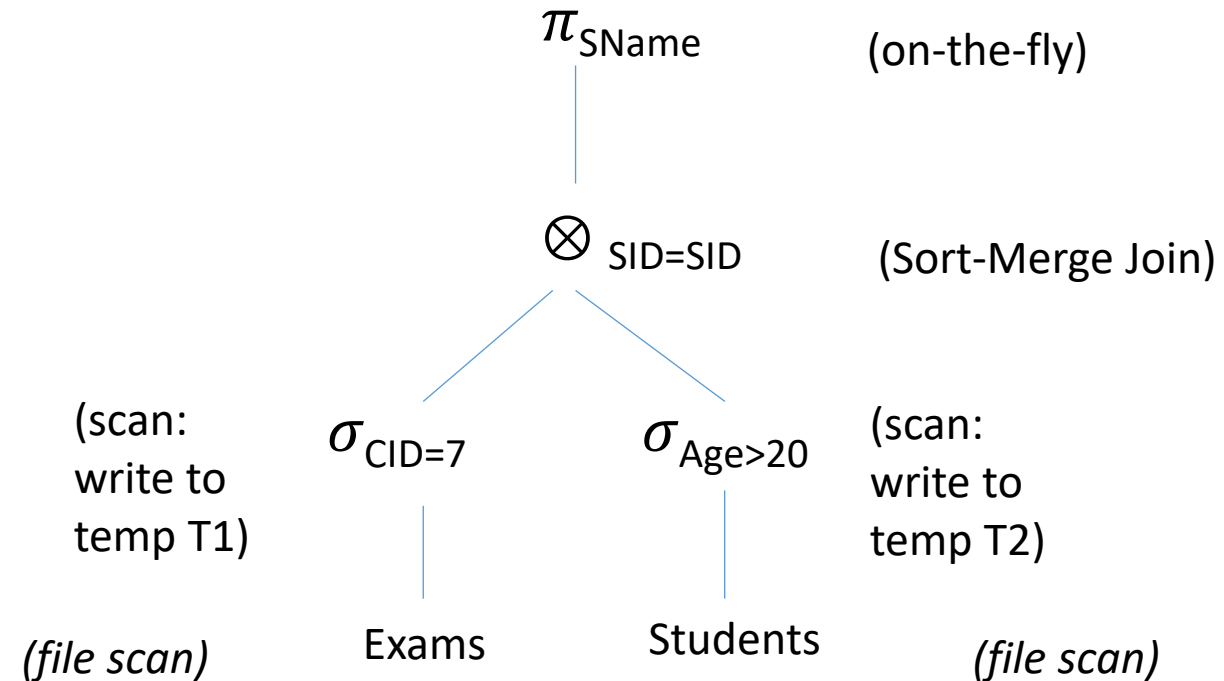
- σ, π – on-the-fly
- cost of plan – very high:
 - $1000 + 1000 * 500 = 501,000$ I/Os



Motivating Example

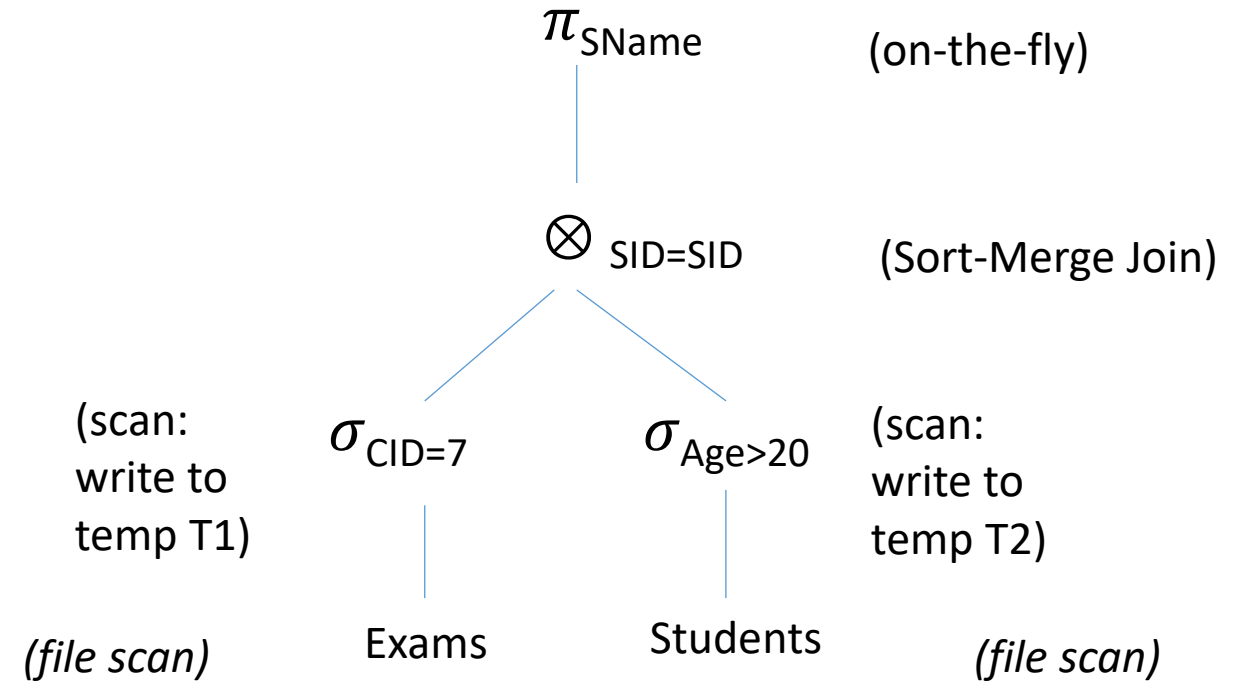
* optimizations

- reduce sizes of the relations to be joined
 - push selections, projections ahead of the join
- alternative plans
 - push selections ahead of joins
- selection
 - file scan
 - write the result to a temporary relation on disk
- join the temporary relations using Sort-Merge Join



Motivating Example

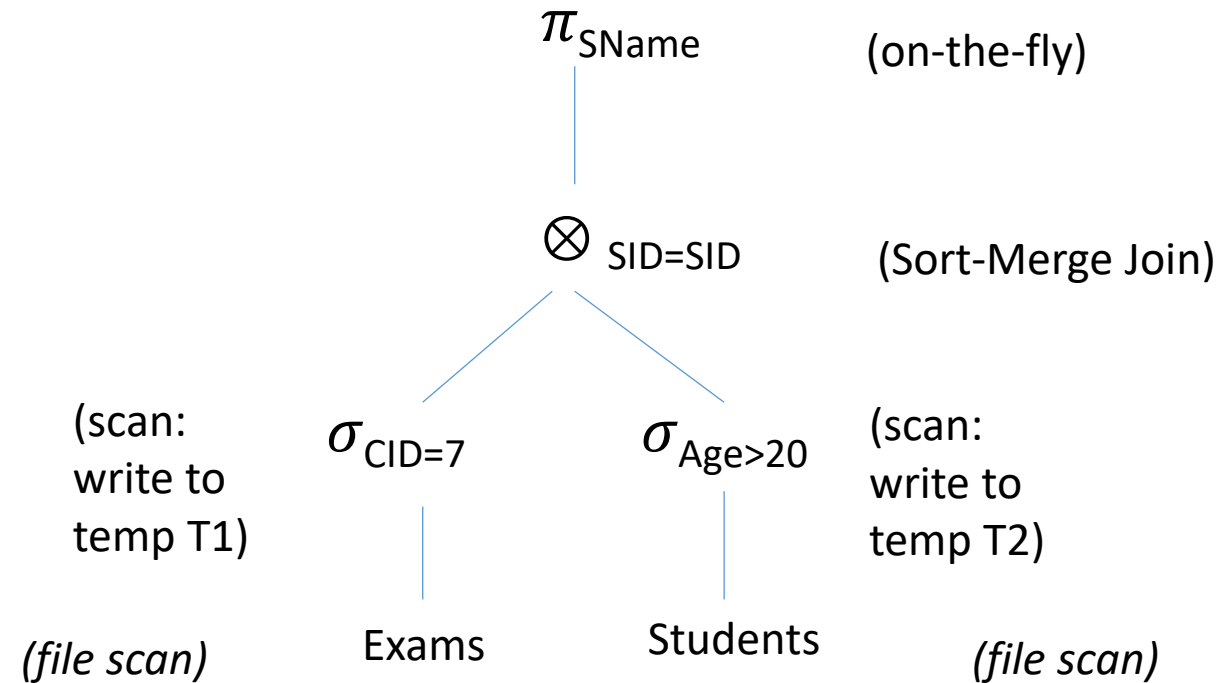
- 5 available buffer pages
- cost
 - $\sigma_{CID=7}$
 - scan Exams: 1000 I/Os
 - write T1
 - assume exams are uniformly distributed across all courses, i.e., T1 has 10 pages (there are 100 courses)
 - $\sigma_{Age>20}$
 - scan Students: 500 I/Os
 - write T2
 - assume ages are uniformly distributed over the range 19 to 22, i.e., T2 has 250 pages



Motivating Example

- 5 available buffer pages
- cost
 - Sort-Merge Join
 - T1 - 10 pages
 - sort T1: $2 * 2 * 10 = 40$ I/Os
 - T2 - 250 pages
 - sort T2: $2 * 4 * 250 = 2000$ I/Os
 - merge sorted T1 and T2
 - $10 + 250 = 260$ I/Os
 - π - on the fly

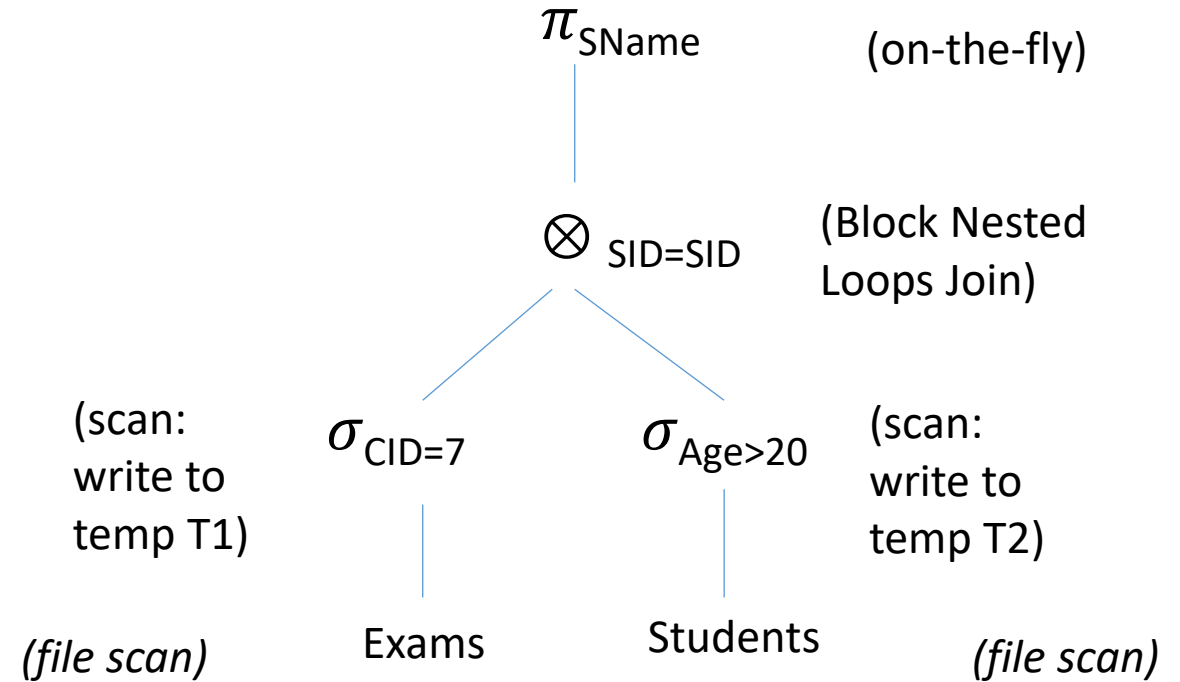
=> **total cost:** $\underbrace{1000 + 10 + 500 + 250}_{\text{selection}} + \underbrace{40 + 2000 + 260}_{\text{join}} = 4060 \text{ I/Os}$



Motivating Example

- 5 available buffer pages
- cost
 - Block Nested Loops Join
 - T1 - 10 pages, T2 - 250 pages
 - T1 - outer relation
=> scan T1: 10 I/Os
 - $\lceil 10/3 \rceil = 4$ T1 blocks
=> T2 scanned 4 times: $4 * 250 = 1000$ I/Os
 - BNLJ cost: $10 + 1000 = 1010$ I/Os
 - π - on the fly

=> **total cost:** $\underbrace{1000 + 10 + 500 + 250}_{\text{selection}} + \underbrace{10 + 1000}_{\text{join}} = \mathbf{2770 \text{ I/Os}}$



Motivating Example

- push projections ahead of joins
 - drop unwanted columns while scanning Exams and Students to evaluate selections => T1[SID], T2[SID, SName]
- T1 fits within 3 buffer pages
 - => T2 scanned only once
 - => **total cost**: about **2000 I/Os**

References

- [Ra02] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (3rd Edition), McGraw-Hill, 2002
- [Da03] DATE, C.J., An Introduction to Database Systems (8th Edition), Addison-Wesley, 2003
- [Ga09] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book (2nd Edition), Pearson, 2009
- [Ra07] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, McGraw-Hill, 2007,
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si19] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts (7th Edition), McGraw-Hill, 2019
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems,
<http://infolab.stanford.edu/~ullman/fcdb.html>